

Introduction to TCP and its behaviour on the network

Toke Høiland-Jørgensen - Red Hat

Agenda


- Connection startup
- Loss recovery mechanisms
- Congestion control
- Interactions with other parts of the network
- Measurement tools and techniques

TCP conceptually


- Connection oriented
 - Have to connect () before transmitting
- Byte stream based
 - Chopping data into packets is a protocol detail
- Reliable
 - Data will (eventually) arrive in order, or error out
- Adaptable
 - Will attempt to transmit data as fast as possible over any link
- Old
 - [RFC675](#) was first published 51 years ago - lots has happened since!

Connection startup

TCP handshake



TCP handshake



TCP handshake




Image source: [https://en.wikipedia.org/wiki/Handshake_\(computing\)](https://en.wikipedia.org/wiki/Handshake_(computing))

Adding TLS





Image source: https://en.wikipedia.org/wiki/Transport_Layer_Security#TLS_handshake

Important points - connection startup


- Small packets
- Multiple round-trips to setup a connection
- Dominated by RTT
- Can take longer than the data transfer!

TCP loss recovery

Retransmission timeout (RTO)




Fast retransmit (DupACK)




Selective Acknowledgement (SACK and DSACK)

- SACK: [RFC2018](#)
- DSACK: [RFC2883](#)



Tail Loss Probes (TLP)



Shorter timeout, don't reduce window

Recent Acknowledgement (RACK)

- [RFC8985](#)
- Keep a “recent window” based on RTT (`rtt_min/4`)
- Don't trigger loss recovery if OOO packets arrive within this window
- Significantly improves behaviour when the network reorders packets

Important points - loss recovery

- Several different mechanisms in play
- SACK, RACK and TLP improves behaviour
 - Reordering resilience in particular
- Active research and innovation - moving target!

Congestion control

Congestion control determine TCP behaviour. We'll cover:

- Classic (Reno) TCP
- CUBIC
- BBR

[Wikipedia has a list of other algorithms](#)

Congestion signals

All congestion control algorithms react to one or more signals:

- Packet loss (duplicate ACKs)
- Timeout (RTO)
- ECN notification
- RTT increases

“Slow” start in classic (Reno) TCP




Image source: https://commons.wikimedia.org/wiki/File:TCP_Slow-Start_and_Congestion_Avoidance.svg

Many flows never leave slow start!


TCP CUBIC

- [RFC9438](#)
- Default in Linux, MacOSX and Windows 10+
- More aggressive than RENO
- But still fundamentally loss based

Congestion signals (CUBIC)

- Packet loss (duplicate ACKs)
- Timeout (RTO)
- ECN notification
- RTT increases

TCP CUBIC window function



(b) CUBIC window growth function.

TCP BBR

- First published in 2016
- [draft-ietf-ccwg-bbr](#) (WiP) at the IETF
- Based on *modelling link behaviour*
- Used extensively at Google (e.g., YouTube)
- Multiple versions (v1, v2, v3), only v1 in Linux

Congestion signals (BBRv1)

- Packet loss (duplicate ACKs)
- **Timeout (RTO)**
- ECN notification
- **RTT increases**

Congestion signals (BBRv3)

- Packet loss (duplicate ACKs)
- Timeout (RTO)
- ECN notification
- RTT increases

TCP BBR behaviour




Image source: <https://queue.acm.org/detail.cfm?id=3022184>


Important points - congestion control

- “Classic” (New) Reno often implicit assumption in literature
- CUBIC actually default, and is more aggressive
- BBR totally different, mainly deployed in the cloud (Google)


TCP interactions

Latency

Latency per Bandwidth




Page Load Time As RTT Decreases



Source: <https://youtu.be/TNBkxA313kk> (from 2011)


Bufferbloat



Bufferbloat (cont)

- Loss (or ECN mark) is an important signal for TCP!
- Timely signalling required or TCP will keep scaling up
- Active Queue Management (AQM) and Fairness Queueing (FQ) helps
- Excessive link-layer retries add latency

Head of line (HOL) blocking



Also happens when 2 is reordered, not lost.

TCP hardware offload

We assume traffic is smooth, but:

- Servers use TCP Segmentation Offload (TSO)
- Network stack creates 64KB superpacket, NIC segments
- These go out at line rate (10-100G) as one burst!
- Can overwhelm smaller links downstream

TCP Pacing

- Linux feature that spreads out packets over an RTT
- *Reduces* burstiness by spacing packets
- Helps reduce buffering
- Can hurt aggregation at wireless links

Important points - TCP interactions

- TCP interacts with other layers in non-obvious ways
- Latency equally important as throughput
- Congestion signalling is important!
 - TCP will fill buffers, leading to bloat
- Hardware offloads and pacing can affect packet sequences

Measurement tools and techniques

Basic throughput tests

- iperf (2/3)
- netperf

```
$ sysctl -a | grep congestion
net.ipv4.tcp_allowed_congestion_control = reno cubic bbr
net.ipv4.tcp_available_congestion_control = reno cubic bbr
net.ipv4.tcp_congestion_control = bbr
```

netperf basics - 5 second download test, CUBIC

```
$ netperf -H netperf-eu.bufferbloat.net -t TCP_MAERTS -D -0.2 -l 5 -- -K cubic
MIGRATED TCP MAERTS TEST from 0.0.0.0 (0.0.0.0) port 0 AF_INET to demo.tohojo.dk () port 0 AF_INET : demo
Interim result: 2.81 10^6bits/s over 0.212 seconds ending at 1761825840.513
Interim result: 15.50 10^6bits/s over 0.202 seconds ending at 1761825840.715
Interim result: 33.73 10^6bits/s over 0.200 seconds ending at 1761825840.916
Interim result: 36.40 10^6bits/s over 0.201 seconds ending at 1761825841.117
Interim result: 29.82 10^6bits/s over 0.201 seconds ending at 1761825841.318
Interim result: 24.52 10^6bits/s over 0.201 seconds ending at 1761825841.519
Interim result: 25.66 10^6bits/s over 0.200 seconds ending at 1761825841.720
Interim result: 26.21 10^6bits/s over 0.201 seconds ending at 1761825841.920
Interim result: 25.90 10^6bits/s over 0.200 seconds ending at 1761825842.121
Interim result: 27.07 10^6bits/s over 0.202 seconds ending at 1761825842.323
Interim result: 27.15 10^6bits/s over 0.200 seconds ending at 1761825842.523
Interim result: 27.28 10^6bits/s over 0.201 seconds ending at 1761825842.724
Interim result: 28.69 10^6bits/s over 0.201 seconds ending at 1761825842.926
Interim result: 25.47 10^6bits/s over 0.202 seconds ending at 1761825843.127
Interim result: 29.19 10^6bits/s over 0.200 seconds ending at 1761825843.327
Interim result: 27.83 10^6bits/s over 0.204 seconds ending at 1761825843.531
Interim result: 29.02 10^6bits/s over 0.202 seconds ending at 1761825843.733
Interim result: 27.22 10^6bits/s over 0.200 seconds ending at 1761825843.933
Interim result: 28.79 10^6bits/s over 0.200 seconds ending at 1761825844.133
Interim result: 29.09 10^6bits/s over 0.200 seconds ending at 1761825844.333
Interim result: 30.31 10^6bits/s over 0.201 seconds ending at 1761825844.534
Interim result: 27.07 10^6bits/s over 0.204 seconds ending at 1761825844.738
Interim result: 28.90 10^6bits/s over 0.200 seconds ending at 1761825844.938
Interim result: 28.04 10^6bits/s over 0.202 seconds ending at 1761825845.140
Interim result: 29.53 10^6bits/s over 0.161 seconds ending at 1761825845.301
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  10^6bits/sec

131072 16384 16384 5.00 26.77
```

netperf basics - 5 second upload test, BBR

```
$ netperf -H netperf-eu.bufferbloat.net -t TCP_STREAM -D -0.2 -l 5 -- -K bbr
MIGRATED TCP STREAM TEST from 0.0.0.0 (0.0.0.0) port 0 AF_INET to demo.tohojo.dk () port 0 AF_INET : demo
Interim result: 13.42 10^6bits/s over 0.215 seconds ending at 1761826036.765
Interim result: 99.22 10^6bits/s over 0.267 seconds ending at 1761826037.032
Interim result: 26.78 10^6bits/s over 0.421 seconds ending at 1761826037.453
Interim result: 48.11 10^6bits/s over 0.311 seconds ending at 1761826037.763
Interim result: 34.58 10^6bits/s over 0.330 seconds ending at 1761826038.093
Interim result: 36.35 10^6bits/s over 0.310 seconds ending at 1761826038.403
Interim result: 37.26 10^6bits/s over 0.317 seconds ending at 1761826038.720
Interim result: 33.66 10^6bits/s over 0.339 seconds ending at 1761826039.059
Interim result: 38.14 10^6bits/s over 0.309 seconds ending at 1761826039.368
Interim result: 36.31 10^6bits/s over 0.314 seconds ending at 1761826039.682
Interim result: 35.89 10^6bits/s over 0.314 seconds ending at 1761826039.996
Interim result: 36.84 10^6bits/s over 0.310 seconds ending at 1761826040.306
Interim result: 37.22 10^6bits/s over 0.317 seconds ending at 1761826040.623
Interim result: 35.58 10^6bits/s over 0.321 seconds ending at 1761826040.943
Interim result: 37.06 10^6bits/s over 0.304 seconds ending at 1761826041.247
Interim result: 37.20 10^6bits/s over 0.303 seconds ending at 1761826041.550
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  10^6bits/sec

131072 16384 16384 5.61 34.40
```


Measuring latency under load

- ping
- fping
- irtt
- smokeping (for continuous monitoring)
- Passive ping

Application-level tests

- [http-getter](#)
- [oha](#)

TCP flow and state analysis

- [Wireshark](#)
- ss utility
- [TCPtrace](#)

ss TCP stats

```
$ ss -ntip | grep -A 1 netperf
ESTAB 0      0      10.42.18.12:45497      193.10.227.30:12865 users:(("netperf",pid=1593163,fd=3))
      bbr wscale:7,10 rto:265 rtt:64.873/29.682 ato:40 mss:1288 pmtu:1420 rcvmss:656 advmss:1368
      cwnd:11 bytes_sent:656 bytes_acked:657 bytes_received:656 segs_out:4 segs_in:3 data_segs_out:1
      data_segs_in:1 bbr:(bw:209424bps,mrtt:49.089,pacing_gain:2.88672,cwnd_gain:2.88672)
      send 1747168bps lastsnd:1206 lastrcv:1157 lastack:1157 pacing_rate 4386232bps
      delivery_rate 209904bps delivered:2 app_limited busy:49ms reordering:4 rcv_space:13312
      rcv_ssthresh:13312 minrtt:49.089 snd_wnd:64512 rcv_wnd:65536
--
ESTAB 0      3776416 10.42.18.12:45271      193.10.227.30:41287 users:(("netperf",pid=1593163,fd=4))
      bbr wscale:7,10 rto:382 rtt:181.719/6.64 mss:1288 pmtu:1420 rcvmss:536 advmss:1368 cwnd:624
      bytes_sent:4792648 bytes_acked:3988937 segs_out:3723 segs_in:1044 data_segs_out:3721
      bbr:(bw:43716968bps,mrtt:38.541,pacing_gain:2.88672,cwnd_gain:2.88672) send 35382629bps
      lastsnd:1 lastrcv:1119 lastack:1 pacing_rate 124936600bps delivery_rate 33778600bps
      delivered:3098 delivered_ce:2649 busy:1119ms unacked:624 reordering:4 rcv_space:13680
      rcv_ssthresh:64168 notsent:2972704 minrtt:38.541 snd_wnd:3144192 rcv_wnd:65536
```

tcptrace example


```
$ sudo tcpdump -epni wg0 -s 128 -w tcp-trace.pcap host netperf-eu.bufferbloat.net
# run test and exit with Ctrl-C afterwards
tcpdump: listening on wg0, link-type RAW (Raw IP), snapshot length 128 bytes
^C5856 packets captured
5862 packets received by filter
0 packets dropped by kernel

$ tcptrace -G tcp-trace.pcap
1 arg remaining, starting with 'tcp-trace.pcap'
Ostermann's tcptrace -- version 6.6.7 -- Thu Nov  4, 2004

TCP packet 1: reserved bits are not all zero.
    Further warnings disabled, use '-w' for more info
5856 packets seen, 5856 TCP packets traced
elapsed wallclock time: 0:00:00.062305, 93989 pkts/sec analyzed
trace file elapsed time: 0:00:05.517782
TCP connection info:
  1: alrua-x1:52789 - 193.10.227.30:12865 (a2b)    7>    5<  (complete)
  2: alrua-x1:38581 - 193.10.227.30:41197 (c2d) 1298> 4546<  (complete)
$ xplot c2d_tsg.xpl
# interactive viewer
```

wireshark example

```
$ wireshark tcp-trace.png  
Statistics -> TCP Stream Graphs -> Time sequence (tcptrace)
```



Integrated test tools

- Flent
- bbperf

Flent basic example

```
$ flent tcp_2up netperf-eu.bufferbloat.net --socket-stats
Starting Flent 2.2.0 using Python 3.13.7.
Starting tcp_2up test. Expected run time: 70 seconds.
Data file written to ./tcp_2up-2025-10-30T131517.453720.flent.gz
```

Summary of tcp_2up test run from 2025-10-30 13:15:17

	avg	median	99th %	# data	pts
Ping (ms) ICMP :	123.28	142.00	181.28 ms		237
TCP upload avg :	17.69	N/A	N/A Mbits/s		350
TCP upload sum :	35.38	N/A	N/A Mbits/s		350
TCP upload::1 :	19.34	20.23	24.48 Mbits/s		350
TCP upload::1::tcp_bbr_bw :	22.33	22.61	28.29		217
TCP upload::1::tcp_bbr_cwnd_gain :	1.99	2.00	2.89		217
TCP upload::1::tcp_bbr_mrtt :	43.90	42.53	45.97		217
TCP upload::1::tcp_bbr_pacing_gain :	1.03	1.00	2.62		217
TCP upload::1::tcp_cwnd :	300.49	308.00	392.64		218
TCP upload::1::tcp_delivery_rate :	19.26	19.37	23.22		218
TCP upload::1::tcp_pacing_rate :	22.67	22.58	29.48		218
TCP upload::1::tcp_rtt :	144.11	146.85	173.39		217
TCP upload::1::tcp_rtt_var :	3.51	2.64	15.10		217
TCP upload::2 :	16.04	16.35	21.36 Mbits/s		350
TCP upload::2::tcp_bbr_bw :	18.47	18.24	22.39		218
TCP upload::2::tcp_bbr_cwnd_gain :	1.98	2.00	2.89		218
TCP upload::2::tcp_bbr_mrtt :	44.17	43.20	123.49		218
TCP upload::2::tcp_bbr_pacing_gain :	1.04	1.00	2.61		218
TCP upload::2::tcp_cwnd :	249.99	251.00	326.00		218
TCP upload::2::tcp_delivery_rate :	15.99	15.85	20.08		218
TCP upload::2::tcp_pacing_rate :	18.84	18.13	27.71		218
TCP upload::2::tcp_rtt :	143.85	146.78	172.63		217
TCP upload::2::tcp_rtt_var :	3.63	2.99	14.44		217

Flent basic example (cont)


```
$ flent tcp_2up netperf-eu.bufferbloat.net --socket-stats --test-parameter cc_algos=cubic,cubic
Starting Flent 2.2.0 using Python 3.13.7.
Starting tcp_2up test. Expected run time: 70 seconds.
Data file written to ./tcp_2up-2025-10-30T132010.247839.flent.gz
```

Summary of tcp_2up test run from 2025-10-30 13:20:10

	avg	median	99th %	# data	pts
Ping (ms) ICMP :	71.80	72.70	102.04 ms		350
TCP upload avg :	16.34	N/A	N/A Mbits/s		350
TCP upload sum :	32.67	N/A	N/A Mbits/s		350
TCP upload::1 :	13.67	13.89	18.05 Mbits/s		350
TCP upload::1::tcp_cwnd :	104.85	105.00	126.66		218
TCP upload::1::tcp_delivery_rate :	13.10	13.21	17.64		218
TCP upload::1::tcp_pacing_rate :	17.42	17.40	23.25		218
TCP upload::1::tcp_rtt :	75.16	74.52	98.30		218
TCP upload::1::tcp_rtt_var :	3.77	3.04	16.36		218
TCP upload::2 :	19.00	18.98	26.00 Mbits/s		350
TCP upload::2::tcp_cwnd :	146.07	146.00	173.00		218
TCP upload::2::tcp_delivery_rate :	18.09	18.12	23.91		218
TCP upload::2::tcp_pacing_rate :	24.40	23.98	33.08		218
TCP upload::2::tcp_rtt :	74.57	74.59	98.65		218
TCP upload::2::tcp_rtt_var :	3.57	2.78	13.41		218


Flent basic example - plotting throughput and latency

```
# for interactive display
$ flent-gui tcp_2up*.flent.gz
# for saving
$ flent tcp_2up*.flent.gz -f plot --subplot-combine -o tcp-tests.png
```



Flent basic example - plotting TCP window size (from ss)

```
$ flent tcp_2up*.flent.gz -f plot --subplot-combine -p tcp_cwnd -o tcp-tests-cwnd.png
```



BBR probing and CUBIC window function clearly visible!

Speedtests

- [Waveform Bufferbloat test](#)
- [Cloudflare speed test](#)
- [speedtest.net](#)
- [fast.com](#)

Pay attention to latency under load figures!

Important points - measurement

- Basic throughput tests a good first step
 - Netperf / iperf both usable, netperf has more features
- Measure latency as well (and under load!)
- Deep-diving into TCP behaviour can be necessary
- Using integrated measurement tools is helpful
- Speed tests have latency measurements as well now!

Further reading

- Wikipedia: https://en.wikipedia.org/wiki/TCP_congestion_control
- Bufferbloat web site: <https://bufferbloat.net/>
- My own PhD: <https://bufferbloat-and-beyond.net/>
- “It’s the latency, stupid”: <https://www.stuartcheshire.org/rants/Latency.html>
- RFCs can be very readable (links throughout the presentation)

